

Graph Decomposition and Parity

Bobby DeMarco^{1,*} and Amanda Redlich^{2,†}

¹ E-mail: rvdemarco@gmail.com

²DEPARTMENT OF MATHEMATICS
BOWDOIN COLLEGE
BRUNSWICK, MAINE, 04011
E-mail: aredlich@bowdoin.edu

Received February 14, 2013; Revised July 7, 2015

Published online in Wiley Online Library (wileyonlinelibrary.com).
DOI 10.1002/jgt.21907

Abstract: Motivated by a recent extension of the zero-one law by Kolaitis and Kopparty, we study the distribution of the number of copies of a fixed disconnected graph in the random graph $G(n, p)$. We use an idea of graph decompositions to give a sufficient condition for this distribution to tend to uniform modulo q . We determine the asymptotic distribution of all fixed two-component graphs in $G(n, p)$ for all q , and we give infinite families of many-component graphs with a uniform asymptotic distribution for all q . We also prove a negative result that no recursive proof of the simplest form exists for a uniform asymptotic distribution for arbitrary graphs. © 2015

Wiley Periodicals, Inc. *J. Graph Theory* 00: 1–13, 2015

1. INTRODUCTION

A recent paper by Kolaitis and Kopparty [3] gives an extension of the zero-one law, which holds for first-order logic with a parity operator. The keystone of their proof is that the number of copies (not necessarily induced) of any fixed connected graph is

*Contract grant sponsor: U.S. Department of Homeland Security; Contract grant number: 2007-ST-104-000006;

†Contract grant sponsor: National Science Foundation; Contract grant number: 1004382.

asymptotically uniformly distributed modulo q for any q in the random graph $G(n, p)$. Other papers have studied this statistic in special cases. For example, a 2004 paper of Loeb et al. [4] considered the modulo q distribution of triangles, while the 2014 work of Kopparty and Gilmer [1] gives the distribution of triangles overall.

Here we study the distribution of the number of copies (not necessarily induced) of a fixed *disconnected* graph in $G(n, p)$ modulo q . For convenience, we may say “the distribution of a graph modulo q ” to mean the distribution of the number of copies of the graph in $G(n, p)$ with p implicit. In this article, we will only be speaking of asymptotic distributions, so when discussing distributions we will remove the word asymptotic for brevity. We say a graph is uniformly distributed if it is uniformly distributed modulo q for all q . We give sufficient conditions for a graph to be uniformly distributed, and we use these conditions to completely characterize the distribution of all two-component graphs for all q . We then give infinitely large families of uniformly distributed graphs of any component size.

In analyzing these distributions, we developed the concepts of *unique composition* and *decomposition*. These concepts are related to determining when several connected graphs may be combined to create one large connected graph with certain uniqueness properties. There are obvious links to the reconstruction conjecture (see [2] for a summary), which asks when the subgraphs of a larger graph have slightly different uniqueness properties. In this article we give an algorithm for uniquely composing any two feasible graphs, and certain families of three or more graphs. We also show no generic recursive composition algorithm exists.

The article is structured as follows. In the second section, we derive a formula for the number of copies of a disconnected graph in a fixed graph G as a function of the number of copies of certain connected graphs in G . In the third section, we use this formula to give specific conditions for a disconnected graph to be uniformly distributed.

We show these conditions are satisfied for almost all two-component graphs in the fourth section. We give an explicit construction for all satisfying graphs. We also calculate the distribution for all two-component graphs that do not satisfy these conditions. We then give some examples of infinite families of three or more component graphs that are satisfying. We conclude this section with a negative result, showing no simple algorithm exists to show a generic graph is satisfying. In the last section, we discuss areas of further research.

2. COUNTING COPIES

In this section, we give an exact formula for the number of unlabeled copies of a disconnected graph $A = \sqcup_{i=1}^k G_i$, with nonisomorphic G_i in a host graph F . That is, the number of subgraphs $G_A \subseteq F$ such that $A \simeq G_A$. For example, if F is K_4 and A is two disjoint edges, the number of unlabeled copies of A in F is 3. Note that this is different from the labeled case, which would give 24 copies.


Our formula for the number of unlabeled copies of a disconnected graph is given in terms of the number of copies of various connected graphs H , and their relationship to the original graph A . Although the formula appears complex, the reasoning behind it is simple. The main idea is that each copy of A is the product of copies of G_1, G_2, \dots, G_k . Interactions between copies of G_i and G_j lead to an overcount; the formula uses the principle of inclusion and exclusion to correct for this. Notice that this is correct only if the G_i are nonisomorphic, hence our assumption that each component is unique.

The simplest case is when $A = G_1 \sqcup G_2$. For example, let $G_1 = C_3$ and $G_2 = C_4$. Given fixed host graph F , let $N(A)$ be the number of unlabeled copies of A in F (we may also mean the number of unlabeled copies of A in an instance of the random graph $G(n, p)$; this will be clear from context). Consider H_i as illustrated below. For ease of discussion, here and throughout the article, we label the illustrated vertices. However, the graphs themselves are unlabeled.

We have

$$N(A) = N(C_3)N(C_4) - N(H_1) - N(H_2) - 2N(H_3) - 3N(H_4).$$

That is because the total number of disconnected C_3, C_4 pairs is the total number of C_3, C_4 pairs minus the number of connected pairs. While H_1 and H_2 each correspond to exactly one connected pair, H_3 is counted twice; once when $C_3 = \{1, 2, 3\}$ and once when $C_3 = \{1, 3, 4\}$. Similarly, H_4 is counted three times: $C_4 = \{1, 2, 3, 4\}$ or $\{1, 2, 3, 5\}$ or $\{1, 5, 3, 4\}$, and each copy of C_4 determines a complementary copy of C_3 . (Note that we are counting both induced and noninduced subgraphs; for example, H_3 contains no induced C_4).



We now formalize this “gluing” idea.

Definition 1. A tuple $(G_1, G_2, \dots, G_k, H, H_1, H_2, \dots, H_k)$ is a *gluing* of $G_1 \dots G_k$ if

- H is a connected graph
- H_1, \dots, H_k are subgraphs of H
- $H_i \sim G_i$ for all $i \in [k]$
- $\cup_{i=1}^k E(H_i) = E(H)$.

We occasionally refer to H itself as a *gluing*. The tuple (H_1, \dots, H_k) is a *decomposition* of H . If there exists only one tuple H_1, \dots, H_k such that $(G_1, \dots, G_k, H, H_1, \dots, H_k)$ is a gluing, we say that (G_1, \dots, G_k, H) is *uniquely decomposable*. In this case, we say that $H[S]$ is the unique subgraph of H induced by $\{H_i\}_{i \in S}$. We occasionally say that H itself is *uniquely decomposable*, if G_1, \dots, G_k are clear from context. If there exists an H such that (G_1, \dots, G_k, H) is uniquely decomposable, we say that $\{G_1, \dots, G_k\}$ is *uniquely composable*.

We often want to count gluings and decompositions.

Definition 2. Given G_1, \dots, G_k , $s(H)$ is the number of tuples (H_1, \dots, H_k) such that $(G_1, \dots, G_k, H, H_1, \dots, H_k)$ is a gluing. The set of gluings \mathbf{H} is the family of graphs H such that $s(H) \neq 0$.

Using this notation, we restate a theorem often used in the theory of graph limits and implied in [3].

Theorem 3. For a disconnected graph $A = G_1 \sqcup G_2$ and $G_1 \neq G_2$,

$$N(A) = N(G_1)N(G_2) - \sum_{H \in \mathbf{H}} s(H)N(H).$$

Proof. The number of copies of A is the number of G_1, G_2 pairs overall, less the number of G_1, G_2 pairs that intersect. Each intersecting G_1, G_2 pair corresponds to a decomposition of a copy of some H , so the total number of intersecting pairs is $\sum_{H \in \mathbf{H}} s(H)N(H)$. ■

It is tempting to generalize this to the three or more component case as

$$N(A) \text{ “} = \text{”} \prod_{i=1}^k N(G_i) - \sum_{H \in \mathbf{H}} s(H)N(H),$$

but the truth is more complicated. Along with H that may be decomposed into G_1, \dots, G_k , we must consider H that are decomposable into any subset of G_i . For example, if A has components G_1, G_2, G_3 , then we must be concerned with gluings of the forms (G_1, G_2, H, H_1, H_2) , (G_1, G_3, H, H_1, H_3) , (G_2, G_3, H, H_2, H_3) , and $(G_1, G_2, G_3, H, H_1, H_2, H_3)$. In order to deal with this complication, we define some new terms. First, some notation about partitions.

Definition 4. Consider the partitions of $[k]$ under partial ordering by refinement, where we use π and ρ for partitions of $[k]$ and say $\pi < \rho$ if π is a refinement of ρ . If $\pi < \rho$, and T is a block of ρ , then let $\pi(T)$ be the family of blocks in π such that $\cup_{S \in \pi(T)} S = T$. For example, if $k = 4$, $\pi = \{\{12\}\{3\}\{4\}\}$, and $\rho = \{\{123\}\{4\}\}$, then $\pi(\{123\}) = \{\{12\}\{3\}\}$.

Now we use partitions to classify gluings. Each connected component corresponds to a set in a partition, and the family of gluings is broken into subfamilies according to the partitions they generate.

Definition 5. Let \mathbf{H}_π be the family of graphs $H = \sqcup_{S \in \pi} H_S$, where each H_S may be decomposed (not necessarily uniquely) into $\{H_i\}_{i \in S}$. For example, $\mathbf{H}_0 = A$ and any $H \in \mathbf{H}_1$ is a connected graph; $\mathbf{0}$ is the minimum and $\mathbf{1}$ is the maximum element in the partial order.

We also count possible decompositions of gluings. Since we are now considering a broader range of gluings, we must add a subscript to clarify which graphs are being glued.

Definition 6. Given G_1, \dots, G_k and $S \subseteq [k]$, let $s_S(G')$ be the number of ways the graph G' may be decomposed into copies of $\{G_i\}_{i \in S}$.

It may be useful to discuss decompositions into graphs other than our original components G_i .

Definition 7. Given a graph H_S for each $S \in \pi$, $s_{\pi(T)}(G')$ is the number of ways G' may be split into $\{H_S\}_{S \in \pi(T)}$. If there exists some $\{H_S\}_{S \in \pi}$ such that $s_{\pi([k])}(G') \neq 0$, say that G' is compatible with π .

Finally, we need to count “component” decompositions.

Definition 8. For any graph H , let $p(H)$ be the number of isomorphic permutations of the components of H . That is, if the components of H are i_1 copies of some connected graph B_1 , i_2 copies of B_2 , up to i_k copies of B_k , then $p(H) = i_1! i_2! \cdot \dots \cdot i_k!$. For example, if H is a five-component graph consisting of three C_4 and two K_5 , then $p(H) = 3!2!$: there are $3!$ ways to decide which C_4 is which and $2!$ ways to decide which K_5 is which.

This notation allows us to give a more general recursion for the number of copies of a graph with an arbitrary number of connected components. Although the notation is daunting, it is a simple generalization of the ideas in the two-component case.

Theorem 9. For a graph $A = \sqcup_{i=1}^k G_i$ with $G_i \neq G_j$ for all $i \neq j$,

$$\begin{aligned} N(A) &= \prod_{i=1}^k N(G_i) - \sum_{\mathbf{0} < \pi \leq [k]} \sum_{H \in \mathbf{H}_\pi} N(\sqcup_{S \in \pi} H_S) p(H) \prod_{S \in \pi} s_S(H_S) \\ &= \prod_{i=1}^k N(G_i) - \sum_{\mathbf{0} < \pi \leq [k]} \sum_{H \in \mathbf{H}_\pi} p(H) \prod_{S \in \pi} s_S(H_S) \\ &\quad \left(\prod_{S \in \pi} N(H_S) - \sum_{\rho > \pi} \sum_{J \in \mathbf{H}_\rho} p(J) \prod_{T \in \rho} s_{\pi(T)}(J_T) p(J_T) N(\sqcup_{T \in \rho} J_T) \right). \end{aligned}$$

Proof. Now that we have the proper definitions, the proof is short. As usual, we count the number of copies of A by finding the product of copies of its components, then subtract the overcount. The “overcounted” graphs are those in which at least two G_i intersect with each other, that is, those corresponding to a non- $\mathbf{0}$ partition.

We also need to be careful about the possibility of intersecting components being isomorphic to a third component. For instance, returning to our initial $C_3 \sqcup C_4$ example, if we let A be the three-component graph $C_3 \sqcup C_4 \sqcup H_1$, one of the gluings in $H_{\{1,2\}\{3\}}$ is $H_1 \sqcup H_1$. This gluing should be counted twice because there are two choices for which component is generated by $C_3 \sqcup C_4$. Therefore we have the first line of the equation.

To see why the second line is true, simply apply the first equation to each $N(\sqcup_{S \in \pi} H_S)$ term individually. Now the relevant partitions are those of which π is a refinement, and the decompositions are not into G_i but instead H_S . ■

This theorem gives a recursive algorithm for calculating $N(A)$ for an arbitrary host graph. With sufficient computing power, then, we could use it to calculate $N(A)$ for the random graph directly. This theorem can be used as a starting point that will allow us to give explicit counts of a family of graphs, as well as a sufficient condition for graphs to have certain distributions. The first step is to expand the recursion to get a simpler formula.

Lemma 10. For any graph $A = \sqcup_{i=1}^k G_i$ with $G_i \neq G_j$ for all $i \neq j$, there exist integers $f_A(H)$ for every $H \in \cup_{\pi \leq [k]} \mathbf{H}_\pi$ such that

$$N(A) = \prod_{i=1}^k N(G_i) - \sum_{\pi \leq [k]} \sum_{H \in \mathbf{H}_\pi} \prod_{S \in \pi} N(H_S) f_A(H).$$

Note that f_A is uniquely determined; there is no way to write the number of copies of any connected graph in terms of the number of copies of other connected graphs. This is clear by inspection, or from [3]’s proof that the copies of distinct connected graphs are independently distributed.

3. DISTRIBUTION OF COPIES

As mentioned in the introduction, [3] proves that, for any constants p and q , any $i < q$, and any connected graph G_0 , the probability of $G(n, p)$ having i copies of G_0 modulo q tends to $1/q$ as n tends to infinity. That is, the distribution of a connected graph in the random graph tends to uniform modulo q . We give exact distributions for the number of copies of any disconnected G_0 in $G(n, p)$ modulo q in this section by combining the formulas of the previous section with these results on connected graphs.

The previous section gives exact expressions for the number of copies of a disconnected graph in a particular graph. The formulas are often difficult to implement. However, since our goal is the distribution of the count, rather than its exact value, the preceding formulas are enough. To study the distributions of disconnected graphs, we first recall Theorem 3.2 in [3], which we restate here.

Theorem 11. *For any $q > 1$ and $p \in (0, 1)$, and any family of distinct finite connected graphs F_1, \dots, F_l , the distribution of $(N(F_1), \dots, N(F_l))$ modulo q is $2^{-\Omega(n)}$ close to uniform over $[q]^l$.*

In other words, there are no relations between the number of copies of *distinct* connected subgraphs. Therefore we may use [9] without being concerned about possible dependencies in the distribution. Combining Theorem 11 and Lemma 10 produces the following corollary (note that for any connected graph, e.g., $H \in \mathbf{H}_1$, $p(H) = 1$).

Corollary 12. *Given $A = \sqcup_{i=1}^k G_i$ with distinct G_i , if there exists some $H \in \mathbf{H}_1$ such that $f_A(H)$ is relatively prime to q and H is not a gluing of $\{G_i\}_{i \in S}$ for any $S \subset [k]$, then the distribution of $N(A)$ in $G(n, p)$ is $2^{-\Omega(n)}$ close to uniform modulo q .*

The rest of this section is concerned with finding, for a given A , an H such that $f_A(H) = \pm 1$. This is enough to show that, for any q , $N(A)$ is distributed uniformly modulo q .

Definition 13. *Given $A = \sqcup_{i=1}^k G_i$ and a uniquely decomposable gluing H with decomposition (H_1, \dots, H_k) , the structure graph of H , denoted by $T(H)$, is the graph whose vertices are $[k]$ and edges are pairs i, j such that $H_i \cap H_j$ is nonempty.*

Definition 14. *Given $A = \sqcup_{i=1}^k G_i$, a gluing H is tree-like if it is uniquely decomposable and its structure graph, $T(H)$, is a tree.*

We now show the following.

Theorem 15. *For any graph $A = \sqcup_{i=1}^k G_i$ with distinct components and tree-like gluing H , $f_A(H) = (-1)^{k-1}$.*

Proof. The proof is by strong induction. When $k = 2$ the statement follows from Theorem 3: when $k = 2$ any uniquely decomposable gluing is tree-like.

Now consider $k \geq 3$. By the first line of Theorem 9,

$$N(A) = \prod_{i=1}^k N(G_i) - \sum_{0 < \pi \leq [k]} \sum_{J \in \mathbf{H}_\pi} N(\sqcup_{S \in \pi} J_S) \prod_{S \in \pi} s_S(J_S).$$

The induction hypothesis applied to $\sqcup_{S \in \pi} J_S$ implies that

$$f_A(H) = - \sum_{0 < \pi \leq [k]} \sum_{J \in \mathbf{H}_\pi} f_J(H) \prod_{S \in \pi} s_S(J_S). \quad (1)$$

Note that for any J with $f_J(H) \neq 0$, the unique decomposability of H gives that there is exactly one $0 < \pi \leq [k]$ such that $J \in \mathbf{H}_\pi$. Let us call this partition $\pi(J, H)$. Furthermore, it also gives that $p(J) = 1$ for any such J . We will show the following. ■

Claim 16. *For any r -component J , if $f_J(H) \neq 0$, then $J_S = H[S]$ for all $S \in \pi(J, H)$ and $f_J(H) = -1^{r-1}$.*

Claim 17. *If $f_J(H) \neq 0$ then $\prod_{S \in \pi(J, H)} s_S(J_S) = 1$.*

Claim 18. *For $0 < \pi \leq [k]$, the number of $J \in \mathbf{H}_\pi$ such that $J_S = H[S] \forall S \in \pi$ is one if H is compatible with π and zero otherwise. If H is compatible with π , call π H -good.*

Claim 19. *There are $\binom{k-1}{r-1}$ H -good partitions π consisting of r sets.*

Combining Claims 16–19 with (1) we immediately have (think of r as the number of components of J , or equivalently number of sets in the partition π related to J)

$$\begin{aligned} f_A(H) &= - \sum_{r=1}^{k-1} \binom{k-1}{r-1} (-1)^{r-1} \\ &= -1^{k-1}. \end{aligned}$$

Proof of Claim 16. The first part of the claim holds because H is uniquely decomposable. Note that H is a tree-like gluing of the components of J . Thus the second part of the claim is an application of our inductive hypothesis. ■

Proof of Claim 17. This follows from the unique decomposability of H . ■

Proof of Claim 18. J_S must be a connected graph. So if $H[S]$ is connected for all $S \in \pi$ then $J = \cup_{S \in \pi} H[S]$ is clearly the only $J \in \mathbf{H}_\pi$ such that $J_S = H[S] \forall S \in \pi$. If $H[S]$ is disconnected, then H was not compatible with π . Thus there are zero such graphs. ■

Proof of Claim 19. Consider the natural mapping from a partition $0 < \pi \leq [k]$ to the set $E(T(H) \setminus F)$, where $F = \cup_{S \in \pi} H[S]$. This mapping defines a bijection from the H -good partitions π consisting of r distinct sets and the set of subgraphs of $T(H)$ with $r-1$ edges. ■

4. SPECIFIC EXAMPLES

Here we give some applications of the theorems of the previous section. We begin with a complete characterization of the distribution of all two-component graphs, together with explicit constructions. We then give several families of graphs that have tree-like gluings, and therefore by Theorem 15 are uniformly distributed. Finally, we show that no recursive proof of the simplest form exists for a uniform asymptotic distribution for arbitrary graphs.

4.1. Two Component Graphs

Any uniquely decomposable gluing of two graphs must be tree-like. So one way to show that $N(A)$ is uniformly distributed for some two component graph A would be to give

a construction of a uniquely decomposable H . In fact, such a construction exists for all two-component graphs *except* a few trivial cases.

Theorem 20. *If $G_1 \neq G_2$, neither G_1 nor G_2 is a single vertex, $\{G_1, G_2\} \neq \{P_1, P_2\}$ and $\{G_1, G_2\} \neq \{P_1, P_3\}$, there exists a graph H such that (G_1, G_2, H, H_1, H_2) is a uniquely decomposable gluing and $H \neq G_1, G_2$. Furthermore, H may be constructed explicitly.*

In order to describe the construction of H , we define a few new terms. H will be created by taking two graphs and “gluing” them together.

Definition 21. *Given G_1 and G_2 and vertices $v_1 \in V(G_1)$ and $v_2 \in V(G_2)$, to glue v_1 and v_2 to create a new graph H_v is the natural identification:*

$$V(H_v) = (V(G_1) \setminus \{v_1\}) \cup (V(G_2) \setminus \{v_2\}) \cup \{v\}$$

and

$$\begin{aligned} E(H_v) = & \{\{x, y\} \mid \{x, y\} \in E(G_1) \cup E(G_2)\} \cup \{\{x, v\} \mid \{x, v_1\} \\ & \in E(G_1)\} \cup \{\{v, y\} \mid \{v_2, y\} \in E(G_2)\}. \end{aligned}$$

Given G_1 and G_2 and edges $\{u_1, v_1\} = e_1 \in G_1$ and $\{u_2, v_2\} = e_2 \in G_2$, to glue e_1 and e_2 to create a new graph H_e is the natural identification:

$$V(H_e) = (V(G_1) \setminus \{u_1, v_1\}) \cup (V(G_2) \setminus \{u_2, v_2\}) \cup \{u, v\}$$

and

$$\begin{aligned} E(H_v) = & \{\{x, y\} \mid \{x, y\} \in E(G_1) \cup E(G_2)\} \cup \{\{x, u\} \mid \{x, u_1\} \in E(G_1) \cup \{x, u\} \mid \{x, u_2\} \\ & \in E(G_2)\} \cup \{\{x, v\} \mid \{x, v_1\} \in E(G_1)\} \cup \{\{x, v_2\} \mid \{x, v_2\} \in E(G_2)\} \cup \{\{u, v\}\}. \end{aligned}$$

Notice that there are two possible gluings along an edge, as there are two possible pairings of the endpoints of the edges. Constructions in this article work for an arbitrary pairing of endpoints.

Another difficulty is deciding where to glue two graphs. To describe gluing locations, we consider the underlying connectivity structure of each graph. Recall that a “block” in a graph is a maximal two-connected subgraph; the block structure of a graph concerns relationships between these subgraphs. Inspired by this viewpoint, we say the *block degree* $b_G(v)$ of a vertex $v \in G$ is the number of components generated by the removal of v , i.e. $b(v) = (\text{number of components of } G - v) - (\text{number of components of } G)$. Note that $b_G(v) > 0$ if and only if v is a cut vertex. So every connected graph has at least two vertices of block degree 0, which we will call *block-leaves*. Let $B(G) = \max_{v \in V(G)} b_G(v)$.

Throughout the following discussion we let H refer to the graph created by gluing together G_1 and G_2 at either v_1 and v_2 , or e_1 and e_2 , as discussed in Definition 21. H_1, H_2 will be an arbitrary decomposition of H . That is, H_i may be the original graph G_i , or it may be a different image of G_i in H . Note that, if H is formed by gluing at a vertex, then $H_1 \cap H_2$ is a single vertex. Similarly, if H is formed by gluing at an edge, $H_1 \cap H_2$ is a single edge. We begin with a few observations about the block degree.

Observation 22. *If H is made by gluing together G_1 and G_2 at a vertex, then in any decomposition H_1, H_2 , with $z = H_1 \cap H_2$, for all $x \neq z \in H_1$,*

$$b_{H_1}(x) = b_H(x)$$

and for all $x \neq z \in H_2$,

$$b_{H_2}(x) = b_H(x).$$

Furthermore, $b_{H_1}(z), b_{H_2}(z) \leq b_H(z)$.

Observation 23. If H is formed by gluing together G_1 and G_2 at an edge, then in any decomposition H_1, H_2 , for all $x \in H_1 \cap \bar{H}_2$,

$$b_{H_1}(x) = b_H(x)$$

and for all $x \in H_2 \cap \bar{H}_1$,

$$b_{H_2}(x) = b_H(x).$$

For $x \in H_1 \cap H_2$, $b_{H_1}(x), b_{H_2}(x) \leq b_H(x)$.

With these definitions and observations in hand, we begin the proof of Theorem 20.

Proof. Without loss of generality, we may assume $B(G_2) \geq B(G_1)$ and, if $B(G_2) = B(G_1)$, then $|V(G_1)| \leq |V(G_2)|$. Let S_i be the set of vertices in G_i of block-degree $B(G_i)$. We split graph pairs into six cases, according to their block degrees and other traits, and give a construction for each case.

Case A: $B(G_2) > B(G_1)$ and $B(G_2) > 1$. In this case, glue a block-leaf at maximum distance from S_2 to any block-leaf in G_1 to create H .

Suppose $H_2 \cap G_1 \neq \{v\}$. $H_2 \cap G_1$ is a connected graph: if not, then because H_2 is connected there is a path between any two disconnected components of $H_2 \cap G_1$ within G_2 . But any such path must begin and end at v , and therefore $H_2 \cap G_1$ itself was connected.

Thus there are at least two block-leaf vertices in $H_2 \cap G_1$, hence at least one block-leaf in $H_2 \cap G_1$ not equal to v . Choose one such vertex and label it w . Let R be the set of vertices in H_2 such that $\phi(R) = S_2$. Since $b_H(v) = 2 < B(G_2)$, by Observation 1, we must have $R = S_2$. Then $d(w, R) = d(w, S_2) > d(v, S_2)$ a contradiction. Thus $H_2 \cap G_1 = \{v\}$, and the decomposition is unique.

Case B: $B(G_2) = B(G_1)$. In this case, glue a block-leaf in G_2 to any vertex in S_1 . Since $b_H(v) > B(G_2)$, Observation 1 implies that $H_1 \cap H_2 = \{v\}$. Therefore each component of $H \setminus \{v\}$ must be entirely contained within H_1 or H_2 . Now we use that $|V(G_2)| \geq |V(G_1)|$ to conclude H is uniquely decomposable.

Case C: $B(G_2) = 1, B(G_1) = 0, G_1 \neq K_2$, and there exists a block within G_2 that is not isomorphic to G_1 . Because $B(G_2) = 1$, each cut vertex in G_2 connects two blocks. Therefore we can look at G_2 's structure as a tree T_{G_2} , with vertices $v \in T_{G_2}$ corresponding to each block in G_2 , and edges $e \in T_{G_2}$ corresponding to each cut vertex $u \in G_2$.

Color a vertex in this tree black if its block in G_2 is isomorphic to G_1 and white otherwise. Let a special path in T_{G_2} be any path v_1, v_2, \dots, v_k in T_{G_2} such that v_1 is white and v_2, \dots, v_k are black. Notice that by definition there must be a white vertex, therefore there must be at least one special path (possibly consisting of just one vertex).

Let u_1, u_2, \dots, u_M be a special path of maximal length. Create H by gluing a vertex of G_1 to a vertex of u_M not in u_{M-1} . H now has a longer special path than G_2 . This is uniquely decomposable because any decomposition of H must break this new longest special path by splitting a single vertex. The only vertex that will split H into a copy of G_1 and a copy of G_2 is the glued vertex; any other will generate two graphs, each of which contains at least two blocks.

Case D: $B(G_2) = 1, B(G_1) = 0, G_1 \neq K_2$, and G_2 consists of blocks isomorphic to G_1 .

Glue any edge of e_1 of G_1 to any edge in e_2 of G_2 whose end vertices both have block degree 0 (such an edge must exist, as G_2 consists of blocks of G_1 , and all edges in G_1 have end vertices of block degree 0). Call this new glued block $G_g \subset H$, and the glued edge e . We show that H is uniquely decomposable by focusing on G_g and how it must interact with any decomposition H_1, H_2 .

First notice that G_g cannot be entirely contained within H_1 : G_g has too many vertices. Furthermore, G_g cannot be entirely contained within H_2 : G_g is not isomorphic to G_1 , so it cannot be a block in H_2 . Therefore G_g contains the intersection of H_1 and H_2 . We also know that H_1 is entirely within G_g : If H_1 has vertices both in and out of G_g , it would contain a cut vertex. Also notice a simple counting shows that the intersection of H_1 and H_2 must consist of a single edge.

Therefore, we have established that $H_1 \subset G_g$, $G_g \cap H_2 \neq \emptyset$, and $H_1 \cap H_2 = e'$ for some edge $e' \in G_g$. We now show that $e' = e$.

Suppose not. Notice that, by construction, there are no edges other than e in H that have one end vertex in G_1 and the other end vertex in G_2 . Therefore, if $e' \neq e$, e' must be entirely within G_1 or entirely within G_2 . Thus splitting along e' means splitting G_1 or G_2 . If G_1 is split, one of H_1, H_2 is a proper subgraph of G_1 , which is a contradiction. If G_2 is split, then G_1 is a proper subgraph of both H_1 and H_2 , which is a contradiction.

Case E: $B(G_2) = 1$, $D(G_2) > 2$ and $G_1 = K_2$. In this case, if G_2 contains a vertex of degree one, glue a vertex of G_1 to a leaf at maximum distance from S'_2 , the set of vertices of G_2 of maximum degree. Let w be the vertex in G_1 not glued to G_2 . Note that $d(w, S'_2) = d(v, S'_2) + 1$, which is strictly greater than the distance from x to S'_2 for any leaf $x \in G_2$. Thus $w \notin H_2$ and we conclude the decomposition of H is unique.

If G_2 does not contain any vertices of degree 1, then glue any vertex of G_1 to any vertex of G_2 . Then w , the vertex in G_1 not glued to G_2 , must be in H_1 and we conclude the decomposition of H is unique.

Case F: $G_2 = P_k$, $k > 3$, and $G_1 = K_2$. In this case, glue a vertex of K_2 to the third vertex along the path P_k . It is clear that this graph is uniquely decomposable. ■

In fact, this construction covers almost all uniformly distributed two-component graphs. We fully characterize the distributions of two-component graphs by combining Theorem 20 with some examination of a few special cases.

Theorem 24. *For every graph A with connected components $G_1 \neq G_2$*

- *If neither G_1 nor G_2 is a single vertex, and $\{G_1, G_2\} \neq \{P_1, P_2\}, \{P_1, P_3\}$ (where P_i is the path with i edges), $N(A)$ is $2^{-\Omega(n)}$ close to uniformly distributed in $G(n, p)$ modulo any q .*
- *If $A = P_1 \sqcup P_2$, $N(A)$ is $2^{-\Omega(n)}$ close to uniformly distributed modulo q*
- *If $A = P_1 \sqcup P_3$, $N(A)$ is $2^{-\Omega(n)}$ close to uniformly distributed modulo q if and only if q is odd. If q is even, $N(A)$ is $2^{-\Omega(n)}$ close to being*

$$P(N(A) \equiv 2i) = 3/2q$$

and

$$P(N(A) \equiv 2i + 1) = 1/2q$$

for all $i \in \{0, \dots, q/2\}$.

- *If, without loss of generality, $G_1 = K_1$, $N(A)$ is $2^{-\Omega(n)}$ close to being*

$$P(N(A) \equiv il) = l/q,$$

where $l = \gcd(q, n - |V(G_2)|)$

Proof. The first item follows directly from Theorem 20, Theorem 15, and Corollary 12. The next cases, $P_1 \sqcup P_2$, $P_1 \sqcup P_3$, and $K_1 \sqcup G_2$ are solved by direct computation.



First consider the case $A_1 = P_1 \sqcup P_2$. All gluings $H \in \mathbf{H}$ of P_1 and P_2 are illustrated above. Note that $H_2 = P_2$, $s(H_1) = 2$, $s(H_2) = 2$, $s(H_3) = 3$, and $s(H_4) = 3$. Therefore we have


$$N(A_1) = (N(P_1) - 2)N(P_2) - 2N(H_1) - 3N(H_3) - 3N(H_4).$$

By Theorem 11, we know that the tuple $(N(P_1), N(P_2), N(H_1), N(H_3), N(H_4))$ is $2^{-\Omega(n)}$ close to being uniformly distributed over \mathbf{Z}_q^5 . Therefore $N(A_1)$ itself is $2^{-\Omega(n)}$ close to being uniformly distributed.

Now consider the case $A_2 = P_1 \sqcup P_3$. All gluings $H \in \mathbf{H}$ of P_1 and P_3 are illustrated below. Note that $H_2 = P_3$, $s(H_1) = 2$, $s(H_2) = 3$, $s(H_3) = 2$, $s(H_4) = 4$, and $s(H_5) = 2$. Therefore we have

$$N(A_2) = (N(P_1) - 3)(N(P_3)) - 2N(H_1) - 2N(H_3) - 4N(H_4) - 2N(H_5).$$

Again, Theorem 11 and some basic modular arithmetic are enough to generate the distributions modulo q in each case.



Now, consider the case $A_3 = K_1 \sqcup G_2$. It is clear that $N(A) = (n - |V(G_2)|)N(G_2)$. Once more, Theorem 11 and some basic modular arithmetic are enough to generate the distribution modulo q . ■

4.2. Tree-Like Gluings

Graphs with more than two components are harder to work with using the methods of the previous section. As the number of components increases, the possible gluings and decompositions also increase. Nevertheless, there are some families of multicomponent graphs that admit a recursive construction.

Theorem 25. *If $A = \sqcup_{i=1}^k G_i$ and there do not exist $i \neq j$ such that G_i is a subgraph of G_j , then there exists H_A a tree-like gluing of $\{G_i\}$ such that $H_A \not\cong G_i$, and $N(A)$ is $2^{-\Omega(n)}$ close to being uniformly distributed modulo q for all q .*

Proof. Without loss of generality, let the graphs be listed in nondecreasing order by diameter. Let u_i and v_i be vertices of G_i at maximal distance from each other. Then let H_A be the graph constructed by gluing v_i to u_{i+1} . The structure graph is clearly a tree.

It is also uniquely decomposable, by induction: Suppose this construction is uniquely decomposable for all $k < n$. Now consider H_A for $k = n$. Suppose there exists some decomposition so that $G_n \not\cong H_n$. Consider v_n . If $v_n \in H_i$ for $i \neq n$, then because $D(G_n) \geq D(G_i)$, all vertices in H_i must be within G_n . That contradicts our initial condition that no graphs is a subgraph of another, so it cannot happen.

Therefore $x \in H_n$. Again, by a diameter argument, $G_n = H_n$. Therefore any decomposition of H_A must fix G_n . The graph $H_A - H_n$ is the construction for G_1, \dots, G_{n-1} , so by induction it is also uniquely decomposable.

Because H_A has a tree structure and is uniquely decomposable, by Theorem 15 $N(A)$ is $2^{-\Omega(n)}$ close to uniformly distributed modulo q .

The reader can generate many corollaries of Theorem 25, using any subgraph-free family of graphs.

The same “path-like” gluing shows another family of multicomponent graphs is also uniquely decomposable.

Theorem 26. *If $A = \sqcup_{i=1}^k G_i$ and the G_i are distinct and two-connected, then $N(A)$ is $2^{-\Omega(n)}$ close to being uniformly distributed modulo q for all q .*

Proof. Similarly to the proof of Theorem 25, glue the graphs together at the vertices of maximum distance from each other. H_A contains $k - 1$ vertices of block degree 1, which are exactly the glued vertices. Notice that any decomposition must split all block degree 1 vertices into two block degree 0 vertices; therefore there is exactly one decomposition and the graph is uniquely decomposable. ■

4.3. No Generic Gluing Exists

The previous constructions used a recursive process to create a uniquely decomposable H for G satisfying certain conditions. A natural goal would be to find a generic recursive process to generate H for arbitrary G . However, no such construction exists.

Theorem 27. *There does not exist a generic recursive construction algorithm C that, for all k and distinct G_1, \dots, G_k , generates a uniquely decomposable H_k . That is, there does not exist an algorithm C that, given G_1, \dots, G_k in that order, constructs uniquely decomposable H_k by first calling C on G_1, \dots, G_{k-1} to generate H_{k-1} , and then calling C on H_{k-1}, G_k .* ■

Proof. Suppose there did exist such a recursive C . Let $C(G_1, \dots, G_{k-1}) = H_{k-1}$. If G_1, \dots, G_{k-1} can be glued together as a proper subgraph of H_{k-1} , then C cannot construct a uniquely decomposable H_k on input $G_1, \dots, G_{k-1}, H_{k-1}$. We note that, for example, $G_1 \subseteq G_2$ is enough to give that G_1, \dots, G_{k-1} can be glued together as a proper subgraph of H_{k-1} .

We also point out that ordering is important to this proof; as far as we know, it is possible that an algorithm exists that, given G_1, \dots, G_k , first analyzes the individual graphs, then calls them in a particular order G_{j_1}, \dots, G_{j_k} . ■

5. OPEN QUESTIONS

There are two main open questions. What disconnected graphs are distributed uniformly? What families of connected graphs are uniquely or tree-like composable?

Theorem 12 gives us one means of studying graph distributions. However, it is not the case that graphs are uniformly distributed exactly when they have tree-like compositions. (Recall that $P_1 \sqcup P_2$ is uniform but is not uniquely composable.) It is possible that a more sophisticated analysis of the formula in Theorem 9 could give a different sufficient condition for graphs to be uniformly distributed.

We have fully characterized the two-component graphs that are uniquely composable, and hence admit tree-like compositions. We believe an approach similar to the two-component construction given here also works for the three-component case. However, increasing the number of components significantly complicates the analysis, and the number of cases is over 20. We are currently developing a simpler construction for three components.

Of course, the ultimate goal is to completely characterize the uniquely composable and tree-like composable graphs with any number of components. We suspect that many, if not all, graphs admit such compositions. Theorem 27 indicates a recursive approach does not work in general, but a different type of algorithm may succeed. Even a nonconstructive proof of the existence of uniquely decomposable or tree-like graphs would be interesting.

REFERENCES

- [1] J. Gilmer and S. Kopparty, A local central limit theorem for the number of triangles in a random graph, [arXiv.org:1412.0257v2](https://arxiv.org/abs/1412.0257v2), 2014.
- [2] F. Harary: A survey of the reconstruction conjecture, In: *Graphs and Combinatorics* (Ruth Bari and Frank Harary, Eds.), *Lecture Notes in Mathematics*, Vol. 406, Springer, Berlin/Heidelberg, 1974, pp. 18–28.
- [3] P. G. Kolaitis and S. Kopparty, Random graphs and the parity quantifier, *J ACM* 60(5) (2013), Article 37.
- [4] M. Loeb, J. Matoušek, and O. Pangrác, Triangles in random graphs, *Discrete Math* 289 (2004), 181–185.